
CS 559 - Fall 2021

Homework #6

Krishna Garg
kgarg8@uic.edu

1 Default Hyperparameters

Optimizer: Adam
Weight Decay: 0
Batch size: 100
Learning rate: 1e-3
Scheduler: StepLR
Scheduler step size: 1
Scheduler gamma: 0.7
Loss: CrossEntropyLoss

2 Network

2-layer CNN: Each layer of Conv2d is followed by Relu activation function, followed by Max Pooling layer & 2 layers each of Dropout and Fully connected layer. The inputs are 3-channel images of shape (200 * 200) and the output is one of the 9 classes corresponding to each image. The training dataset has 72k images and the testing dataset contains 18k images.

3 What worked and what didn't work

3.1 80-20 split with default hyperparameters

Didn't work - 100% training accuracy but 87.40% testing accuracy in 1-2 epochs

3.2 60-20-20 split with default hyperparameters

Didn't work - 100% training accuracy but 87.40% accuracy on both val and test datasets in 1-2 epochs

3.3 Tuning the optimizer Adam, with weight_decay (L2 regularization)

3.3.1 weight_decay = 0.003

Didn't work, gave same results

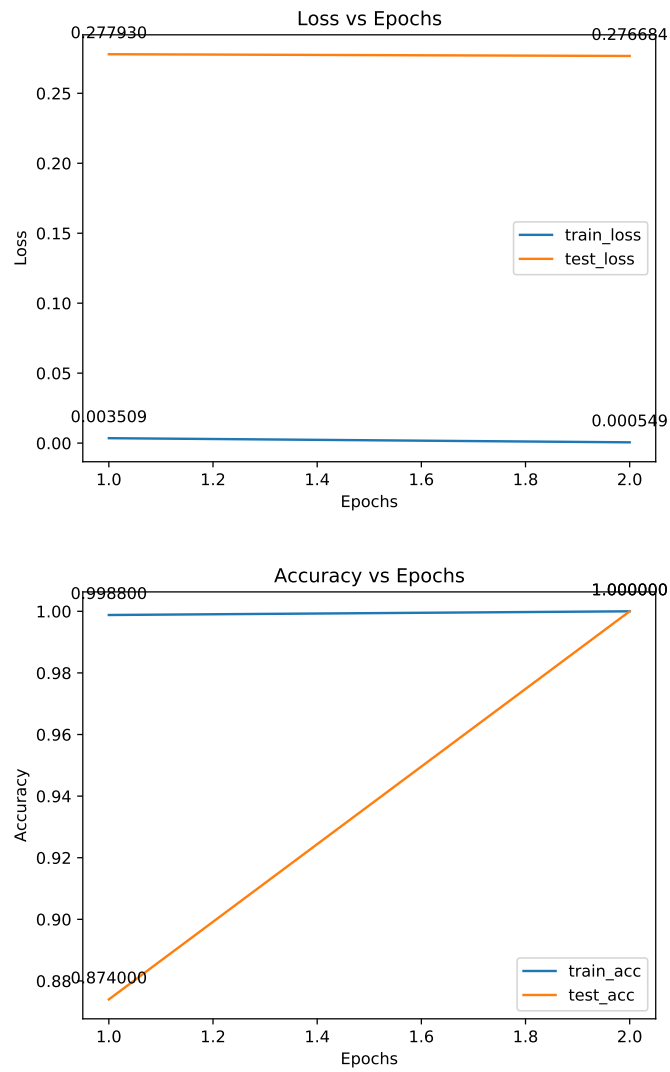


Figure 1: Plots: (a) Loss vs Epochs, (b) Accuracy vs Epochs for both training and testing sets

3.3.2 `weight_decay = 0.03`

Worked, gave 100% accuracy on all datasets. The L2 regularization added helped to overcome the overfitting problem in the above cases and led to build better generalization capability of the model.

Code - 0601-656377418-Garg.py

```
# Code heavily inspired from:
# https://github.com/pytorch/examples/blob/master/mnist/main.py

# training
import argparse, torch, torch.nn as nn, torch.nn.functional as F,
    torch.optim as optim, glob, shutil, random, numpy as np, math
from torchvision import datasets, transforms
from torch.optim.lr_scheduler import StepLR

# Uncomment to create separate folders for train/ test
# random.seed(112)
# for item in ['Circle', 'Square', 'Octagon', 'Heptagon', 'Nonagon',
#             'Star', 'Hexagon', 'Pentagon', 'Triangle']:
#     mylist = [f for f in glob.glob('output/{}*'.format(item))]
#     random.shuffle(mylist)
#     for filename in mylist[:8000]:
#         shutil.copy(filename, 'train_original/images/')
#     for filename in mylist[8000:]:
#         shutil.copy(filename, 'test_original/images/')

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout(0.25)
        self.dropout2 = nn.Dropout(0.5)
        self.fc1 = nn.Linear(614656, 128)
        self.fc2 = nn.Linear(128, 9)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x

def train(args, model, device, train_loader, optimizer, epoch):
    model.train()
    tot_loss = 0
    correct = 0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = torch.nn.CrossEntropyLoss()(output, target)
        loss.backward()
        optimizer.step()

        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

    tot_loss = tot_loss + loss.item()
    if batch_idx % args.log_interval == 0:
        print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f},
              Accuracy: {:.2f}%'.format(
```

```

        epoch, batch_idx * len(data), len(train_loader.dataset),
        100. * batch_idx / len(train_loader),
        tot_loss/(batch_idx+1),
        100.0*correct/((batch_idx+1)*args.batch_size)))

    loss = tot_loss / len(train_loader)
    acc = 100.0 * correct / (len(train_loader) * args.batch_size)
    print('End of Epoch: {}'.format(epoch))
    print('Training Loss: {:.6f}, Training Accuracy:
        {:.2f}%'.format(loss, acc))

def test(args, model, device, test_loader):
    model.eval()
    tot_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            tot_loss += torch.nn.CrossEntropyLoss()(output, target).item()
            # sum up batch loss
            pred = output.argmax(dim=1, keepdim=True) # get the index
                of the max log-probability
            correct += pred.eq(target.view_as(pred)).sum().item()

    loss = tot_loss / len(test_loader)
    acc = 100.0 * correct / (len(test_loader) * args.test_batch_size)
    print('Test Loss: {:.6f}, Test Accuracy: {:.2f}%'.format(loss, acc))
    return loss, acc

def main():
    # Training settings
    parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
    parser.add_argument('--batch-size', type=int, default=100,
        help='input batch size for training (default: 64)')
    parser.add_argument('--test-batch-size', type=int, default=1000,
        help='input batch size for testing (default: 1000)')
    parser.add_argument('--epochs', type=int, default=14, help='number of
        epochs to train (default: 14)')
    parser.add_argument('--lr', type=float, default=1e-3, help='learning
        rate (default: 1.0)')
    parser.add_argument('--gamma', type=float, default=0.7,
        help='Learning rate step gamma (default: 0.7)')
    parser.add_argument('--seed', type=int, default=112, help='random
        seed (default: 112)')
    parser.add_argument('--log-interval', type=int, default=100,
        help='how many batches to wait before logging training status')
    parser.add_argument('--save-model', action='store_true',
        default=True, help='For Saving the current Model')
    args = parser.parse_args()
    print(args.save_model)

    random.seed(args.seed)
    np.random.seed(args.seed)
    torch.manual_seed(args.seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    transform=transforms.Compose([transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))])
    dataset1 = datasets.ImageFolder('train_original/',
        transform=transform)

```

```

dataset2 = datasets.ImageFolder('test_original/',
                                transform=transform)
train_loader = torch.utils.data.DataLoader(dataset1,
                                             batch_size=args.batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset2,
                                             batch_size=args.test_batch_size)

model = Net().to(device)
optimizer = optim.Adam(model.parameters(), lr=args.lr,
                        weight_decay=0.03)

scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
for epoch in range(1, args.epochs + 1):
    train(args, model, device, train_loader, optimizer, epoch)
    test_loss, test_acc = test(args, model, device, test_loader)
    scheduler.step()

    if args.save_model: # save new model after every epoch
        torch.save(model.state_dict(), "0602-656377418-Garg.pt")

    if test_loss < 1e-4 or math.isclose(test_acc, 100.0):
        break

if __name__ == '__main__':
    main()

```

Code - 0603-656377418-Garg.py

```

# inference
import torch, torch.nn as nn, torch.nn.functional as F, random, numpy as np
from torchvision import datasets, transforms

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout(0.25)
        self.dropout2 = nn.Dropout(0.5)
        self.fc1 = nn.Linear(614656, 128)
        self.fc2 = nn.Linear(128, 9)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x

test_batch_size = 100
saved_model_path = '0602-656377418-Garg.pt'
device = torch.device("cuda" if torch.cuda.is_available() else
                      "cpu")
checkpoint = torch.load(saved_model_path, map_location=device)
model = Net().to(device)
model.load_state_dict(checkpoint)

```

```

model.eval()

transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.1307,), (0.3081,))])
test_dataset = datasets.ImageFolder('test_original/',
                                     transform=transform)
test_loader = torch.utils.data.DataLoader(test_dataset,
                                           batch_size=test_batch_size)

tot_loss = 0
correct = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        output = model(data)
        tot_loss += torch.nn.CrossEntropyLoss()(output, target).item() #
            sum up batch loss
        pred = output.argmax(dim=1, keepdim=True) # get the index of
            the max log-probability
        correct += pred.eq(target.view_as(pred)).sum().item()

print('Test Loss: {:.6f}, Test Accuracy: {:.2f}%'.format(
    tot_loss/(len(test_loader)),
    100.0*correct/(len(test_loader)*test_batch_size)))

```
